

**Version: 1.10.0**

Release Notes for EGS version 1.10.0

Release Date: 27th December 2024

The EGS (Elastic GPU Service) platform is an innovative solution designed to optimize GPU utilization and efficiency for your AI projects. EGS leverages the power of Kubernetes to deliver optimized GPU resource management, GPU provisioning and GPU fault identification.

We continue to add new features and enhancements to EGS.

These release notes describe the new changes and enhancements in this version.

What's New

GPU Nodes Idle Timeout

The EGS platform now allows you to configure the idle timeout for GPU nodes. The idle timeout specifies the duration of time that GPU nodes can remain idle before being reallocated to fulfill other GPU requests. Both admins and users can specify the GPU idle timeout during GPR creation.

To know more, see:

- [Create a GPR as an admin](#)
- [Create a GPR as a user](#)

Cost Allocation of Slices

The **Cost Management** that shows cost allocation per slice is now part of the EGS Admin Portal. It provides the KubeTally reports on the utilization of PV, Load Balancer,

CPU, and Memory, providing granular reporting for each slice workspace. However, the cost allocation of GPU nodes are unavailable.

Support for Brownfield in the Cluster Inventory

The cluster inventory now shows all available nodes that is EGS and non-EGS nodes.

Copyright © Avesha 2024. All rights reserved. [Terms and Conditions](#) and [Privacy Policy](#)